

Learning Concept Importance Using a Weighted Dependence Model



Michael Bendersky

Donald Metzler

W. Bruce Croft

UMass Amherst

Yahoo! Research

UMass Amherst



Machine Learning & Friends Lunch, UMass
February, 2010

Input *free-text user query*

The screenshot shows the Yahoo! search interface. At the top, the Yahoo! logo is on the left, and navigation links for Web, Images, Video, Local, Shopping, and More are in the center. A search bar contains the text 'civil war battle reenactments', and a yellow 'Search' button is to its right. Below the search bar, a dropdown menu is visible. On the left side, there is a 'Search Pad' section with a 'Search Scan - On' indicator and a result count of '742,000 results for civil war battle ree...:'. Below this are links for 'Show All' and 'Wikipedia'. The main search results area on the right features a section titled 'Also try: georgia civil war battle reenactments, More...'. Below this, three search results are listed, each with a blue title, a snippet, and a green URL with a 'Cached' link. The first result is 'Annual Gettysburg Civil War Battle Reenactment' with a snippet about the July 2008 event and the URL 'www.gettysburgreenactment.com'. The second result is 'American Civil War reenactment - Wikipedia, the free encyclopedia' with a snippet about the Battle of Chancellorsville and the URL 'en.wikipedia.org/wiki/American_Civil_War_reenactment'. The third result is 'Civil War Reenactors Units, Campaigners' with a snippet about battle schedules and the URL 'www.sutler.net/eventlist.asp'.

YAHOO!

Web Images Video Local Shopping More

civil war battle reenactments Search Options

Search Pad

Search Scan - On

742,000 results for civil war battle ree...:

Show All

Wikipedia

Also try: [georgia civil war battle reenactments](#), [More...](#)

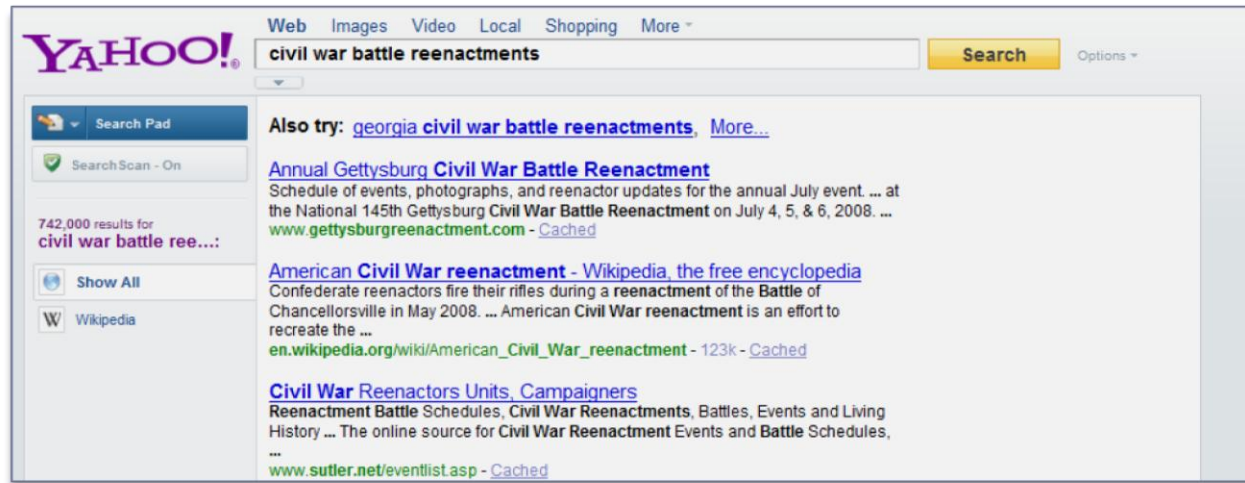
Annual Gettysburg Civil War Battle Reenactment
Schedule of events, photographs, and reenactor updates for the annual July event. ... at the National 145th Gettysburg Civil War Battle Reenactment on July 4, 5, & 6, 2008. ...
[www.gettysburgreenactment.com](#) - [Cached](#)

American Civil War reenactment - Wikipedia, the free encyclopedia
Confederate reenactors fire their rifles during a reenactment of the Battle of Chancellorsville in May 2008. ... American Civil War reenactment is an effort to recreate the ...
[en.wikipedia.org/wiki/American_Civil_War_reenactment](#) - 123k - [Cached](#)

Civil War Reenactors Units, Campaigners
Reenactment Battle Schedules, Civil War Reenactments, Battles, Events and Living History ... The online source for Civil War Reenactment Events and Battle Schedules, ...
[www.sutler.net/eventlist.asp](#) - [Cached](#)

Output *Ranked list of documents*

Ranked Retrieval



- ▶ Search engine must ***accurately*** interpret query intent
 - ▶ Detect phrases
 - ▶ *new york times* \neq *time new york*
 - ▶ Detect relative term/phrase importance
 - ▶ ***CONTINENTAL*** ***airline*** ***BOOKING***



Textual Scoring

- ▶ Generally concerns computing the similarity between two pieces of text
- ▶ This talk will focus on matching (short, textual) **queries** to (long, textual) **documents**
- ▶ Not as popular as some other search problems (e.g., web link analysis), but very important for many search applications



Retrieval Models

1. Query/Document Representation

- ▶ Bag-of-words
- ▶ Bigrams
- ▶ Phrases

2. Query/Document Similarity (Relevance Score)

- ▶ Quantifies how 'relevant' a document is to an information need (expressed by a query)
- ▶ Documents are ***ranked*** by their ***relevance score***



Term-Based Retrieval Models

- ▶ Term-based retrieval models treat the user's query as a “**bag-of-words**”
 - ▶ BM25 (*Robertson et al., 2000*)
 - ▶ Query Likelihood (*Ponte & Croft, 1998*)
 - ▶ DFR (*Amati, 2003*)
- ▶ A simple query model
 - ▶ Term order is interchangeable
 - ▶ Simple collection-based heuristics to weight query terms
 - ▶ e.g., **IDF**
 - ▶ Term weights do not vary based on their context



Concept-Based Retrieval Models

- ▶ Recently, researchers focused on incorporating term dependence into the term-based retrieval models
 - ▶ Markov Random Fields for IR (*Metzler & Croft, 2005*)
 - ▶ BM25 with term proximities (*Song et al., 2008*)
 - ▶ DFR-SD, DFR-FD (*Peng et al., 2007*)
- ▶ A more realistic query model
 - ▶ Term order is important
 - ▶ Captures concepts - dependencies between query terms
- ▶ However, concept weighting is still ***rigid and ad-hoc***
 - ▶ e.g., ***IDF***



Term and Concept Weighting

- ▶ Term and concept weighting are usually handled outside the retrieval model
 - ▶ Automatic removal of redundant query terms
 - ▶ (Kumaran & Carvalho, 2009)
 - ▶ Classifying concepts (noun phrases) into key/non-key classes
 - ▶ (Bendersky & Croft, 2008)
 - ▶ Weighting query terms using regression on their expected performance
 - ▶ (Lease et al., 2009)



Model Desiderata

1. Concept weighting should be integrated into the retrieval model
 - ▶ Avoid pre/post-retrieval processing of queries
2. Retrieval model should handle different types of concepts
 - ▶ Moving beyond the bag-of-words
3. Retrieval model should be optimized to improve ranking
 - ▶ Avoiding metric divergence
4. Retrieval model should be general
 - ▶ Applicable to various document/query types



What is a Concept?

A concept is any syntactic expression that can be matched within a document

- ▶ Not an exhaustive definition
- ▶ Does not capture “semantics”
- ▶ *Practical definition for Information Retrieval*



Examples of Concepts

▶ **Unigrams**

- ▶ (match each of the terms)

... *four white churches. Our own house looked down over the town ...*

▶ **Exact phrases**

- ▶ (match exact phrase)

... *The White House is the official residence and principal workplace of the President of the United States...*

▶ **Proximities**

- ▶ (match unordered phrase in a window of K terms)

...*it was during this construction that the house was painted white ...*




Markov Random Fields for IR

- ▶ In this work we extend the Markov Random Field approach to IR (*Metzler & Croft, 2005*)
 - ▶ provides a general way of modeling a joint distribution
 - ▶ allows incorporating arbitrary scoring functions
 - ▶ easy to train model weights
 - ▶ state-of-the-art performance
 - ▶ TREC Terabyte/Million Query/Web Tracks (2004-2009)




MRFs for IR

- ▶ Encode document and query terms in a graph **G**
 - ▶ vertices represent document/query nodes
 - ▶ edges encode dependence semantics



MRFs for IR (Continued)

- ▶ Potentials over the cliques of **G**
 - ▶ Non-negative functions over clique configurations
 - ▶ Measure query-document similarity
- ▶ Score the document using the joint probability over the cliques of **G**




$$P_{G,\Lambda}(X_1, \dots, X_N) = \frac{1}{Z} \prod_{c \in C(G)} \psi(c; \Lambda)$$




MRFs for IR: Dependence Assumptions


- ▶ MRFs can encode different dependence assumptions



Full Independence
(FI)




Sequential Dependence
(SD)



Full Dependence
(FD)

MRF - Sequential Dependence Model (SD)



- Assume dependence between adjacent terms
- Effectiveness/Efficiency tradeoff
- Empirically proven retrieval performance



SD Ranking Function

- ▶ Associate each clique in the graph with one or more potential function f

$$P(D|Q) \stackrel{rank}{=} \lambda_T \sum_{q \in Q} f_T(q, D) + \lambda_O \sum_{q_i, q_{i+1} \in Q} f_O(q_i, q_{i+1}, D) + \lambda_U \sum_{q_i, q_{i+1} \in Q} f_U(q_i, q_{i+1}, D)$$

how well does q match D ?
[**bag of words** score]

how well does " $q_i q_{i+1}$ " match D ?
[**exact phrase** score]

how well does $\text{prox}(q_i q_{i+1})$ match D ?
[**proximity** score]

Limitations of SD

$$P(D|Q) \stackrel{rank}{=} \lambda_T \sum_{q \in Q} f_T(q, D) + \\ \lambda_O \sum_{q_i, q_{i+1} \in Q} f_O(q_i, q_{i+1}, D) + \\ \lambda_U \sum_{q_i, q_{i+1} \in Q} f_U(q_i, q_{i+1}, D)$$

- ▶ Parameter tying
 - ▶ All matches of the same type are equally weighted
 - ▶ Especially detrimental for verbose queries
- ▶ Instead, we'd like query concept weights to vary



Weighted Sequential Dependence Model (WSD)

- ▶ Allow the parameters to depend on the concept
- ▶ Assume the parameters take a simple parametric form
 - ▶ maintains reasonable model complexity

$$\lambda(q_i) = \sum_{j=1}^{k_u} w_j^u g_j^u(q_i)$$
$$\lambda(q_i, q_{i+1}) = \sum_{j=1}^{k_b} w_j^b g_j^b(q_i, q_{i+1})$$

w - free parameters

g - concept importance features



Defining Concept Importance

- ▶ Features \mathbf{g} define the concept importance

- ▶ Depend on the concept (term/bigram)

$$\lambda(q_i) = \sum_{j=1}^{k_u} w_j^u g_j^u(q_i)$$
$$\lambda(q_i, q_{i+1}) = \sum_{j=1}^{k_b} w_j^b g_j^b(q_i, q_{i+1})$$

- ▶ Independent of a specific document/document corpus
- ▶ Combine several sources for more accurate weighting
 - ▶ **Endogenous Features** – collection dependent features
 - ▶ **Exogenous Features** – collection independent features



Concept Importance Features

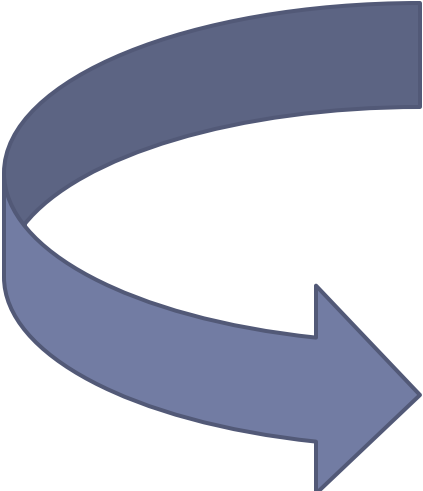
	Data Source	Feature	Description
Endogenous	Collection	$cf(e)$	Collection frequency for concept e
		$df(e)$	Document frequency for concept e
Exogenous	Google n-Grams	$gf(e)$	n -gram count of concept e
	Query Log Sample	$qe_cnt(e)$	# exact query matches for concept e
		$qp_cnt(e)$	# partial query matches for concept e
	Wikipedia Titles	$we_cnt(e)$	# exact title matches for concept e
		$wp_cnt(e)$	# partial title matches for concept e

- **Unigram concepts:** *all features (7)*
- **Bigram concepts:** *all features (7) + PMI for each data source (4)*
- **Total features: 18**
- *All features are log-scaled and normalized*

Incorporating Concept Importance Features

$$\lambda(q_i) = \sum_{j=1}^{k_u} w_j^u g_j^u(q_i)$$

$$\lambda(q_i, q_{i+1}) = \sum_{j=1}^{k_b} w_j^b g_j^b(q_i, q_{i+1})$$


$$P(D|Q) \stackrel{rank}{=} \lambda_T \sum_{q \in Q} f_T(q, D) +$$
$$\lambda_O \sum_{q_i, q_{i+1} \in Q} f_O(q_i, q_{i+1}, D) +$$
$$\lambda_U \sum_{q_i, q_{i+1} \in Q} f_U(q_i, q_{i+1}, D)$$

WSD Ranking Function

- **Score document D by:**

$$\begin{aligned} P(D|Q) \stackrel{rank}{=} & \sum_{i=1}^{k_u} w_i^u \sum_{q \in Q} g_i^u(q) f_T(q, D) + \\ & \sum_{i=1}^{k_b} w_i^b \sum_{q_j, q_{j+1} \in Q} g_i^b(q_j, q_{j+1}) f_O(q_j, q_{j+1}, D) + \\ & \sum_{i=1}^{k_b} w_i^b \sum_{q_j, q_{j+1} \in Q} g_i^b(q_j, q_{j+1}) f_U(q_j, q_{j+1}, D) \end{aligned}$$

- Note that **WSD** model is also linear (with respect to \mathbf{w})



Parameter Estimation

- ▶ Maximum likelihood or maximum *a posteriori* estimation possible from labeled data
- ▶ However, when ranking, it's better to optimize according to some rank-based loss function
- ▶ Many '**learning to rank**' approaches for learning linear ranking models that optimize various retrieval metrics
 - ▶ RankNet, LambdaRank, Ranking SVMs, AdaRank,...



Direct Optimization

- ▶ Learn the weights \mathbf{w} to directly optimize a retrieval performance metric

- ▶ *MAP*
 - ▶ *Mean Avg. Precision*
- ▶ *DCG*
 - ▶ *Discounted Cumulative Gain*

$$P(D|Q) \stackrel{rank}{=} \sum_{i=1}^{k_u} w_i^t \sum_{q \in Q} g_i^u(q) f_T(q, D) + \sum_{i=1}^{k_b} w_i^b \sum_{q_j, q_{j+1} \in Q} g_i^b(q_j, q_{j+1}) f_O(q_j, q_{j+1}, D) + \sum_{i=1}^{k_b} w_i^b \sum_{q_j, q_{j+1} \in Q} g_i^b(q_j, q_{j+1}) f_U(q_j, q_{j+1}, D)$$

- ▶ We use a simple **coordinate-level ascent** algorithm
 - ▶ Efficient for a small number of parameters
 - ▶ Empirically good performance
 - ▶ Most other LR4IR methods can be easily adapted for optimization
-



Query “civil war battle reenactments”

Concept	Importance Features			Weight
	GF	...	DF	
civil	16.9		14.1	<u>0.0619</u>
war	17.9		12.8	<u>0.1947</u>
battle	16.6		12.6	<u>0.0913</u>
reenactments	10.8		9.7	<u>0.3487</u>
civil war	14.5		10.8	<u>0.1959</u>
war battle	9.5		7.4	<u>0.2458</u>
battle reenactments	7.6		4.7	<u>0.0540</u>

Concept weights
may vary even if
concept DF is
similar

Query “civil war battle reenactments”

Concept	Importance Features			Weight
	GF	...	DF	
civil	16.9		14.1	<u>0.0619</u>
war	17.9		12.8	<u>0.1947</u>
battle	16.6		12.6	<u>0.0913</u>
reenactments	10.8		9.7	<u>0.3487</u>
civil war	14.5		10.8	<u>0.1959</u>
war battle	9.5		7.4	<u>0.2458</u>
battle reenactments	7.6		4.7	<u>0.0540</u>


Good segments do not necessarily predict important concepts

Experimental Results


- ▶ A detailed evaluation of our approach
 - ▶ *TREC and web* document collections
 - ▶ *Short & Long* queries
 - ▶ Contribution of *different feature types*
 - ▶ Contribution of *different concept types*



TREC Title (Short) Queries



TREC Description (Long) Queries



Endogenous & Exogenous Features

- ▶ Results with using ***either endogenous or exogenous*** features alone are comparable
- ▶ Using both types of features improves the performance over the unweighted sequential dependence model (SD)
- ▶ In most cases combining both types of features results in better performance



Term & Bigram Weights

- ▶ For short web queries (1-3 terms)
 - ▶ Bigram weights have more impact than term weights
- ▶ For TREC queries and longer web queries
 - ▶ Unigram weights have more impact than bigram weights
- ▶ In most cases combining both types of weights results in better performance, especially for longer queries



Web Queries

	DCG@1	DCG@5	DCG
QL	0.629	1.691	5.844
SD	0.864	2.383	6.681
WSD	0.884 (+2.3%)	2.443 (+2.5%)	6.741 (+0.9%)

- Results using a large-scale commercial web search test collection
- A sample of long web search queries (*length 4+*)
- A total of 1,000 queries with 5-fold CV
- All improvements are statistically significant



Model Desiderata Revisited

1. Concept weighting should be integrated into the retrieval model
 - ▶ ***Concept importance weights are the model parameters***
2. Retrieval model should handle different types of concepts
 - ▶ ***(Potentially) handles any arbitrary term dependencies***
3. Retrieval model should be optimized to improve ranking
 - ▶ ***Concept importance weights are learned to optimize retrieval***
4. Retrieval model should be general
 - ▶ ***Improves retrieval on both newswire & web collections***
 - ▶ ***Improves retrieval with both short & long queries***



Conclusions

- ▶ Existing retrieval methods can be enhanced by
 - ▶ More accurate **modeling** of query concepts
 - ▶ More accurate **weighting** of query concepts
- ▶ Concept weight should be determined by a combination of both endogenous and exogenous features
- ▶ Dynamic concept weighting leads to significant improvements, especially for long queries



Future Work

- ▶ Incorporate concept weighting into the general “learning-to-rank” algorithms
 - ▶ Take into account both textual and non-textual (link-based, click-based) features
- ▶ Extend concept weighting beyond the original query
 - ▶ Query Expansion using (Pseudo-)Relevance Feedback
 - ▶ Query Reformulation
- ▶ Tighter integration of NLP, ML and IR
 - ▶ Using NLP/ML to find meaningful classes of concepts in both queries and documents



Thank you!